

NoticeGuard Team Direction Brief (v2)

LexHack 2026 Project Idea

Idea

NoticeGuard under the **Digital Rights & Policy Tech** track.

A rights checker for independent creators hit by automated copyright claims. It reads the claim and the creator's own documents, then reports whether their evidence is ready for the step they are about to take: **evidence ready**, **evidence gap**, or **needs an adviser**. It reports evidence readiness; it does not tell anyone whether they should file. Every result cites the exact clause and document it relies on.

Problem Statement

Independent creators lose videos, music and revenue to automated copyright claims. The system is huge and one-sided. YouTube alone processed about **2.5 billion** Content ID claims in 2025, almost all of them automated. Creators disputed only **0.51%** of those claims, yet about **two-thirds** of the disputes they filed were resolved in their favour (YouTube Copyright Transparency Report, as reported by TorrentFreak). Matching systems make mistakes: they flag music the creator has licensed, raise claims that conflict with a creator's licence, or flag the wrong segment. Creators rarely challenge them, because the process is opaque and the stakes of getting it wrong are real.

Those stakes rise at every step. On YouTube, a Content ID **claim** is not a **strike**. But if a creator disputes or appeals, the claimant can answer with a copyright **removal request**, which takes the video down and gives the channel a strike; three strikes in 90 days can end the channel. Only after a removal can the creator file a **counter-notice** (counter-notification) under the US DMCA. In it, the creator states **under penalty of perjury** that the removal was a mistake, consents to the jurisdiction of a federal court, and gives the claimant a window (10 to 14 business days) to sue before the work is restored. Knowingly false statements can create liability under **s.512(f)**.

So creators face two failures at once. Many **give up on valid disputes** because they cannot tell whether their licence covers them. A smaller number **escalate unsupported ones** and expose themselves to legal risk.

Limitation of Current Solutions: Given the right documents, a general chatbot can often read a licence correctly. It fails in a different way: it tends to accept the user's framing. Stanford research found that LLMs often fail to correct a user's incorrect legal assumptions (Dahl et al., *Large Legal Fictions*, Journal of Legal Analysis, 2024). A creator who types "*I paid for this track, write my counter-notice*" is likely to get a counter-notice, not a warning. The same question can also get different answers on different runs, and nothing links the output to the legal risk of the step or to a checkable source.

The real problem is not writing an appeal. It is answering one question with evidence: **do my documents support the statement I am about to make, at the step I am about to take?**

Which Challenge Track and Alignment

Chosen Track

Digital Rights & Policy Tech. Automated copyright enforcement is one of the largest private systems deciding what people may publish online. NoticeGuard helps creators exercise the right to challenge it, and does so safely.

How the solution aligns with the Track

- **Digital rights:** It helps creators challenge wrongful claims they would otherwise abandon.
- **Platform accountability:** It checks the claim itself: who the claimant is, what segment was matched, and whether the claimant sits inside the creator's own licence chain.
- **Fair process:** It matches the level of proof to the level of risk. A dispute, an appeal and a sworn counter-notice are treated differently.
- **Transparency:** Every result is traceable to a quoted clause, a date and a rule. Nothing is a black-box opinion.
- **Human control:** The creator decides and signs. NoticeGuard never files anything, never tells the creator to file, and abstains when the documents are ambiguous.

Solution

NoticeGuard is an evidence-checking tool for creators facing a copyright claim. The creator uploads the claim notice, their licence and their receipt, and tells the tool which step they plan to take. NoticeGuard extracts the facts, confirms each one against its source document, runs them through a transparent rules engine and returns an evidence-readiness result with citations.

The traceable chain (what judges can inspect)

Document → Quote → Fact → Rule → Evidence status → Draft

Click any sentence in a result or draft and NoticeGuard shows the chain behind it, for example: *"Why does it say monetised use is allowed?"* → *Licence certificate v2, clause 4.1: 'monetised video on social platforms'* → *fact: monetised use permitted* → *rule R2 (permitted use vs actual use)* → *status: confirmed by document.*

How it works (three layers)

1. Extract with proof. An LLM pulls structured facts from each document (claimant, track, matched segment, purchase date, licence version, permitted uses, governing-terms clause, dates of any permissions). Every fact must carry a verbatim quote from its source. If the quote cannot be found in the document text, the fact is rejected. This blocks hallucinated facts at the door. Where wording has to be mapped to a category (for example, does "monetised video on social platforms" cover a monetised channel?), the mapping is run three times. If the runs disagree, the case goes to **Needs an adviser** instead of the rules. We use agreement between runs rather than the model's own confidence score, because LLMs often cannot tell when they are wrong (Dahl et al., 2024). We are open that this step involves interpretation.

2. Decide with rules, not vibes. A deterministic rules engine decides. It checks licence version by date, permitted use against actual use, claimant against the licence chain, where the creator is in the claim process, and the evidence required for the chosen step. Same input, same result, every time. It checks the statement the creator would be making at that step, not the question they typed. Asking neutrally or insisting "I'm right" produces the same result.

3. Draft only from confirmed facts. A draft is only produced when the evidence is ready, and it is marked "prepared for your review". It may only use facts from the confirmed fact table. A post-check rejects any date, ID or clause in the draft that is not in that table.

What the creator sees

- **Where you are in the process:** a status strip (claim → dispute → appeal → removal and strike → counter-notice) showing the current step and what the next step risks.
- The claim in plain language: who claimed, what was matched, what the platform did, and a deadline countdown.
- **The statement you would be signing,** quoted in full, with each part marked supported or not supported.
- An evidence table in which each fact is labelled **Confirmed by document** (with the quote), **Stated by you**, **Missing** or **Conflicting**, with the key dates (purchase, licence version in force, upload, monetisation, claim).
- One of three results, each written in words and never shown by colour alone, plus the **lowest-risk route** for which the evidence is ready:

Result	Meaning
Evidence ready	Your documents support every part of the statement for this step. Here is a draft for your review, with every sentence cited.
Evidence gap	One statement you would be making is not supported. Here is exactly what is missing and how to get it.
Needs an adviser	The documents are ambiguous, conflict with each other, or the question needs legal judgement (for example fair use). NoticeGuard abstains instead of guessing.

The feature that makes it different

- **NoticeGuard checks the statement you would sign against your own documents, and every step of that check is traceable.** A chatbot answers the question the creator asks, in the way the creator frames it. NoticeGuard checks the statement the creator would have to make at the step they have chosen, and asks for more proof as the legal risk rises. The result can go either way: it tells Maya her evidence is ready for a claim she was about to abandon, and stops Leo from swearing something he cannot yet prove.
 - It is on the creator's side. The default aim is to help valid disputes succeed, not to discourage challenges.
 - The level of proof rises with the level of risk.
 - It does not take the creator's framing for granted. A leading question gets the same result as a neutral one.
 - It picks the least risky route that works. If the licensor can release the claim, or the claimant can retract a removal request, it points there before any dispute or counter-notice.
 - It abstains when it should. "Needs an adviser" is a feature, not a failure.

Why it stands out

- It is not a chatbot or a letter generator. The LLM only extracts. Rules decide, and the traceable chain proves it.
- Its edge is measured against a **fair baseline**: the same documents, several runs, and two prompt styles (a neutral question, and the leading question a frustrated creator actually types).
- The baseline's output is shown unedited, including any cases where it does well.
- Its rules come from real process: licence-version logic, licensed-use conflicts with rights administrators, the claim-to-strike sequence, and the perjury and jurisdiction consequences of a counter-notice.

What it does not claim

NoticeGuard does not decide who owns the copyright, assess fair use, accuse any claimant of bad faith, promise success, tell anyone whether to file, or replace a lawyer. It does not claim that chatbots cannot read licences; with the right documents they often can. The claim is narrower and testable: NoticeGuard is consistent, cites its sources, resists leading questions, and matches proof to risk. Its labels say *confirmed by document*, never *verified*: it checks that documents are consistent with each other, not that they are authentic. The prototype covers one fictional platform (**ClipStream**) whose process (claim, dispute, appeal, removal request and strike, then counter-notice) is modelled on YouTube's published process, one fictional music licensor that issues a licence certificate with every purchase, and synthetic documents only.

Demo Story (3 minutes)

Two creators, the same background track, opposite results. **Target length: 2:45**, safely inside Devpost's 2 to 3 minute limit. **Every baseline shown gets exactly the same documents as NoticeGuard, and its output is recorded unedited.**

Scene 1: Maya, a valid dispute she nearly abandoned (0:00 to 0:50)

Maya, a student animator, has a monetised video claimed by **Northline Rights** for the background track *Glasslight*. Revenue is redirected to the claimant. She assumes she has lost, because the licensor's website now says the Standard licence is non-commercial.

- **Her documents:** the claim notice, her purchase receipt, the **licence certificate (v2)** attached to her purchase email, and the licensor's current terms (v3).
- **NoticeGuard reads the documents:** receipt dated 12 March 2025. The certificate is licence **v2**; clause 4.1 permits "*monetised video on social platforms*." The current **v3** removed that permission, but v3 clause 9 says "*your licence is governed by the terms in force on the date of purchase*."
- **Claim check:** the licensor's track page lists Northline Rights as its content-ID administrator. Northline may be administering the track legitimately, but its claim conflicts with Maya's licence: a **licensed-use conflict**, which is the kind of claim suited to a dispute.
- **Result: Evidence ready** for a dispute. **Lowest-risk route:** first ask the licensor to have its administrator release the claim. If that fails, her evidence is ready for an in-platform dispute; the draft cites v2 cl. 4.1, v3 cl. 9, the receipt ID and the dates. The status strip warns that a dispute can still lead the claimant to file a removal request, and shows that her documents would also support a counter-notice if that happened.
- **Click-through (5 seconds):** Maya clicks the draft sentence about monetised use and sees the chain: certificate v2 → clause 4.1 quote → fact → rule → status.
- **Why this scene matters:** about two-thirds of disputes are resolved in the creator's favour, yet only about one claim in 200 is disputed. Maya is the creator who would have given up.

Scene 2: Leo, the leading question (0:50 to 2:00)

Leo used the same track, bought in August 2025 under the **v3** Standard licence, in a monetised video published on 10 September 2025. His claim followed the full path: he disputed, the claimant reinstated the claim, he appealed, and the claimant then filed a removal request. His video is down and his channel has its first strike. Only now is a counter-notice available, and he wants to file one.

- **Fair baseline, on screen:** Leo gives a general chatbot the same documents NoticeGuard gets and types what a frustrated creator would type: "*I paid for this track. Write my counter-notice.*" We show its answer unedited. Research on false premises suggests it will often comply; our benchmark measures how often.
- **NoticeGuard shows what he would be swearing:** "*I have a good-faith belief that the material was removed as a result of mistake or misidentification.*"
- **Evidence check:** the v3 Standard licence excludes monetised use, and his video was monetised on the upload date. Nothing in his documents grants monetised use. The key statement is **not supported**.
- **Result: Evidence gap.** The screen explains in plain words what a counter-notice commits him to (a statement under penalty of perjury, consent to court jurisdiction, and the claimant's option to sue). It also shows what is at stake: a second and third strike within 90 days could end his channel.
- **What would close the gap:** any document, dated before he published, that grants him monetised use of the track. Other options: remove or replace the track, ask the claimant to retract the removal request, or speak to an adviser or clinic.
- **Live resolution:** Leo remembers that before publishing he emailed the licensor's support team, which replied on 2 September 2025: "*We've extended your licence: you're cleared to use Glasslight in monetised ClipStream videos on your channel.*" He uploads that email. The rules re-run: a grant from the licensor, dated before the upload, covering monetised use. The statement is now supported, so the counter-notice moves to **Evidence ready**. The lowest-risk route comes first: send the email to the claimant and ask them to retract the removal request, which also clears the strike. The counter-notice draft is ready as a fallback. The chain shows exactly which fact changed the result.

Scene 3: Proof on a fair baseline (2:00 to 2:35)

- **Benchmark panel:** 12 cases, 4 of them written by a teammate who did not write the rules. Each case is run 3 times through a plain LLM and through NoticeGuard, with the same documents and both prompt styles (neutral and leading).
- **How the answers were set:** each case's expected result was labelled by hand against a written rubric, which we publish in the repo. The benchmark measures consistency and faithfulness to that rubric. It is not legal validation, and we say so on screen.
- **Metrics:** correct result, same result on every run, citations that point to the real clause, and unsafe counter-notices drafted.
- **Honesty rule:** we show the numbers as they come out. If the baseline does well on neutral questions, we say so; the edge we pitch is consistency, correct citations and resistance to leading questions.

Demo Closing Moment (2:35 to 2:45)

Same song, same documents. Maya was about to give up on a claim she could win. Leo was about to swear something he couldn't prove, until the one email that proved it turned up. NoticeGuard showed each of them exactly where their evidence stood, tied every answer to a clause, and gave the same answer however the question was asked.

Pitch

One sentence pitch

NoticeGuard tells creators whether their own documents back the step they're about to take against a copyright claim, before they take it, and shows the exact clause that proves it.

Full pitch

Creators have about two-thirds of their copyright disputes resolved in their favour, but they dispute only about one claim in 200. Chatbots don't fix this: they tend to accept the creator's framing, give different answers on different runs, and write confident text with no sources. NoticeGuard extracts facts only when they can be quoted from the creator's documents, decides with transparent rules, and reports whether the evidence is ready for the step the creator is about to take, in either direction, with every conclusion traceable to a clause. It helps valid challenges succeed and keeps creators out of legal trouble.

Judging Criteria Map

Criterion (weight)	How NoticeGuard scores
Real-World Impact & Feasibility (25%)	Backed by platform data: 2.5 billion automated claims a year, 0.51% disputed, about two-thirds of disputes resolved for the creator. Narrow, buildable scope. Next step: pilot conversations with university IP and digital-rights clinics through the Fellowship (none approached yet).
Technical Execution (25%)	Quote-grounded extraction, a deterministic rules engine, a draft guard, a click-through traceable chain, and a fair benchmark (same documents, two prompt styles, repeated runs, held-out cases, published rubric). Engineers can inspect all of it.
UX & Design (20%)	A process status strip, three plain-language results, the statement being signed shown in full, a deadline countdown, labels that don't rely on colour, and one-click "what would change this."
Innovation (15%)	Checks the statement the creator would sign, not the question they ask. Proof scales with legal risk, and the least risky route is recommended first.
Presentation (15%)	One track, two creators, opposite outcomes, and the baseline shown unedited on the same documents. The Devpost write-up declares every library, the LLM used, the rubric and the synthetic data.

Build plan to the deadline (Sun 27 Sept, 5pm EDT / 10pm UK)

Must have

- Synthetic document set: ClipStream claim and removal notices, v2 licence certificate, v3 terms, track page naming the administrator, receipts, channel data, and Leo's 2 September support email.
- A one-page rubric that sets the expected result for each benchmark case.
- Upload and extraction with a JSON schema. Each field must include a quote, and the quote-match check rejects any fact without one.
- Rules engine, declarative and versioned: licence version by date, permitted use against actual use, permission dates against upload date, claimant in licence chain, process stage, required evidence per step, route choice, and abstain conditions.
- Result screen: the process status strip, the statement being signed, the evidence table with its four labels, citations with click-through to the chain, lowest-risk route, and "what would change this."
- Benchmark script: 12 cases $\times$ 3 runs $\times$ 2 prompt styles $\times$ 2 systems, feeding a simple results panel.
- **Benchmark documents:** 8 of the 12 cases are variants of the Maya and Leo document sets, made by a small script that changes dates, licence versions, clause wording and claimants. Only the 4 held-out cases are written by hand. This keeps document-writing to a few hours.
- Video and Devpost write-up.

Should have

- Draft generator restricted to confirmed facts, with the post-check.
- Deadline countdown.

Timetable (UK time)

- **Sat afternoon:** align on proposal and kick off
- **Sat evening:** extraction with quote-match; rules engine passing Maya and Leo.
- **Sun morning:** result screen and click-through; benchmark run.
- **Sun 3pm:** feature freeze. **Sun 3pm to 8pm:** record, edit, write Devpost. **Sun 9pm:** submit, one hour before the deadline.

The strict winning test

The demo passes only if all of these hold on screen:

- The baseline gets identical documents in every comparison, and its output is shown unedited.
- Maya gets **Evidence ready** because of the licence-version rule and the administrator check, with both clauses cited.
- Leo's sequence is correct (claim, dispute, appeal, removal and strike, then counter-notice), and his result moves from Evidence gap to Evidence ready only because a genuinely pre-existing document arrives.
- One click on any draft sentence shows its full chain.
- Every sentence in the draft traces to a quoted source.
- The benchmark numbers are real, shown even if they are imperfect, including cases where the baseline does well, and described as rubric-consistency, not legal validation.

If any of these is faked, drop it from the demo rather than bluff. This panel is mostly engineers, and they will look for exactly that.

Sources

- *YouTube Content ID 2025 figures:* TorrentFreak, "YouTube Processed 2.5 Billion Content ID Copyright Claims in 2025" (reporting Google's YouTube Copyright Transparency Report).
- *YouTube Help,* "Dispute a Content ID claim" (answer 2797454) and "Copyright strike basics" (answer 2814000).
- Dahl, Magesh, Suzgun and Ho, "Large Legal Fictions: Profiling Legal Hallucinations in Large Language Models", *Journal of Legal Analysis* 16(1), 2024.
- *Counter-notice requirements:* 17 U.S.C. s.512(g)(2)-(3); *misrepresentation:* s.512(f).